



Microsoft

Azure Machine Learning

Srinag



Principles of Machine Learning

Things to know



Why Machine Learning is the Future?


Measurement of Data

1 TB = 1000 GB


1 PB = 1000 TB

1 EB = 1000 PB

Growth of Data



Statistics on Data Growth



90% of today's data has been created in last two years alone

Statistics on Data Growth




Biotech



Healthcare



Automotive



Internet of Things



Manufacturing



Telecom



Banking and Finance



Social Media



Ecommerce

Heard on the Streets


- IDC Futurescape - Two-thirds of Global 2000 Enterprises CEOs will centre their corporate strategy on digital transformation including machine learning (ML) solutions.
- Harvard Business Review - Data Scientist: The Sexiest Job of the 21st Century
- McKinsey Report - 45 percent of work activities could potentially be automated by currently demonstrated technologies; machine learning can be an enabling technology for the automation of 80 percent of those activities.
- Microsoft CEO Satya Nadella - called out machine learning and the big data that powers it as a key development in his memo to Microsoft last July.




Insights is a journey




Convergence accelerates digital transformation



Cloud




Data





Intelligence

Information Management




Data

Compose and orchestrate data services at scale




- Create, schedule, orchestrate, and manage data pipelines
- Visualize data lineage
- Connect to on-premises and cloud data sources
- Monitor data pipeline health
- Automate cloud resource management
- Move relational data for Hadoop processing
- Transform with Hive, pig, or custom code

Big Data Stores



Data

A hyper-scale repository for big data analytics workloads



- A Hadoop Distributed File System for the cloud
- No fixed limits on file size
- No fixed limits on account size
- Unstructured and structured data in their native format
- Massive throughput to increase analytic performance
- High durability, availability, and reliability
- Azure Active Directory access control

Elastic data warehouse as a service with enterprise-class features

Big Data Stores




Data Lake Store



SQL Data Warehouse




Cosmos DB




- Petabyte scale with massively parallel processing
- Independent scaling of compute and storage—in seconds
- Transact-SQL queries across relational and non-relational data

- Full enterprise-class SQL Server experience
- Works seamlessly with Power BI, Machine Learning, HDInsight, and Data Factory


Machine Learning and Analytics



Easily build, deploy, and share predictive analytics solutions




A screenshot of the Cortana Analytics Gallery interface. The top navigation bar includes "Browse all", "Solution Templates", and "Experiments". The left sidebar has "Refine by" sections for "CATEGORIES" (Solution Template, Experiment, Machine Learning API, Tutorial, Collection, Notebook) and "SHOW" (Microsoft content only). The main area shows "Results" for "MACHINE LEARNING API" and "Face APIs". The "Face APIs" card features an image of two people, a description of Microsoft's state-of-the-art cloud-based face algorithms, and stats (1071687 views, 7 months ago).



- Simple, scalable, cutting edge. A fully managed cloud service that enables you to easily build, deploy, and share predictive analytics solutions.
- Deploy in minutes. Azure Machine Learning means business. You can deploy your model into production as a web service that can be called from any device, anywhere and that can use any data source.
- Publish, share, monetize. Share your solution with the world in the Gallery or on the Azure Marketplace.

Intelligence



Build applications that understand people

Intelligence	Vision	Speech	Language	Knowledge	Search
Cognitive Services	Computer Vision	Speaker Recognition	Text Analytics	Academic	Bing Web Search API
	Content Moderator	Bing Speech	Bing Spell Check	Entity Linking	Bing Image Search API
	Face	Custom Speech Service	Web Language Model	Knowledge Exploration	Bing Video Search API
	Emotion		Linguistic Analysis	Recommendations	Bing News Search API
	Video		Language Understanding	QnA Maker	Bing Auto Suggest API
			Translator		

- Faces, images, emotion recognition and video intelligence
- Spoken language processing, speaker recognition, custom speech recognition
- Natural language processing, sentiment and topics analysis, spelling errors
- Complex tasks processing, knowledge exploration, intelligent recommendations
- Bing engine capabilities for Web, Autosuggest, Image, Video and News

Your bots – wherever your users converse

Intelligence



Cognitive Services



Bot Service



Cortana

Accelerate development cycles

nodeJS

C#

Develop
your way



Built in
code editor



Quick start
templates



Integrated
chat window

Enrich your bots



Channel
support



Cognitive
Services

API

Direct
Line support



Embedded
web chat

Boost operational efficiencies



Powered by
Azure Functions



Continuous
deployment



Scale
on demand



Reduced
dev ops

- Start quickly with built-in templates
- Reach your customers on multiple channels
- Boost the power of bots with Azure services

- No server management or patching needed
- Scale out automatically
- Pay only for what you use

Get things done in more helpful, proactive and natural ways

Intelligence

Cognitive Services

Bot Service

Cortana



Here are some of the things I can help you with...

Answers

Predictions

Monitoring & Alerts

Task Completion

Cortana for Consumers (today)

Public reference data answers – *"How far is it from Los Angeles to San Francisco?"*

Event predictions – *"Who do you think is going to win the Germany Italy game?"*

Flight status, traffic conditions, changes in weather, ...

Setting reminders, scheduling meetings, getting directions, ...

With the Cortana Intelligence

Answers from organizational data in Power BI
"What were our biggest deals that closed last month?"

Integration with prediction solutions
"Which of our customers are most likely to churn in the next quarter?"

Monitoring KPIs and preemptive alerting
"Alert me if this customer ever has a 90% chance of churn in the next 30 days"

Line of business process integration
Assistance with expense report submission on-time within policy

Dashboards & Visualizations




Power BI




- Analytics for everyone, even non-data experts
- Your whole business on one dashboard
- Create stunning, interactive reports
- Drive consistent analysis across your organization
- Embed visuals in your applications
- Get real-time alerts when things change

Transform data into intelligent action





Cloud AI Stack



Machine Learning in Real World


Machine Learning During BREXIT

- The pound sterling fell in value by 11%
- Dow Jones Industrial Average and the London Stock Exchange's FTSE 100 index lost more than 2%
- Omega Point's customers were the least affected
- Reason - Data-driven strategy fuelled by machine learning models



Linkedin

- "People You May Know" ads achieved a click-through rate 30% higher
- They generated millions of new page views.
- Thanks to this one feature, LinkedIn's growth trajectory shifted significantly upward.



Few More Machine Learning Examples



Customers who viewed this item also viewed these products



Weber One Touch Gold
Premium Charcoal
Grill-57cm

\$225

Add to cart

NoMU Salt Pepper and
Spice Grinders

\$3

View options

Hi. I'm Cortana.
Ask me a question!

NETFLIX

Movie Recommendations

HERSHEY'S®

NET WT
5 LB
(2.26 kg)

MILK CHOCOLATE

Fascinating Case Study

- Hershey used IoT sensors and Microsoft Azure algorithms for machine learning to improve production efficiencies on a Twizzler candy line.
- There are 22 sensors on each Twizzler holding tank, with 60 million data points collected.
- Each 1% change in sizing for Twizzlers in a pound holding tank resulted in a savings of \$500,000.
- "We were able to utilize the precooked algorithms inside of Azure to wire up all of the machine learning. We literally were able to build this without a data scientist " said George Lenhart, senior manager, advanced productivity and collaboration, at Industry of Things World USA in San Diego, Calif.



Why Azure ML?

- Drag and Drop interface and no Programming required
- Large variety of algorithm as modules
- From experiment to production API in minutes
- Supports R and Python to bring in your existing code
- Flexibility of data storage; supports variety of data storage options
- Large number of pre-built APIs available as a service

TATA
MOTORS



Rolls-Royce



Benefits of Machine Learning

- Faster decisions
- Develop insights that are beyond human capabilities
- Act at the right time and take advantage of opportunities, converting them into closed deals.


What is Machine Learning?

What is Machine Learning?


- Machine learning is the subfield of computer science that gives computers the ability to learn without being explicitly programmed.
 - Arthur Samuel, 1959
- Learns from past behaviour and make predictions or decisions
- Extraction of knowledge from data

Machine Learning


- Machine learning is the subfield of computer science that gives “computers the ability to learn without being explicitly programmed”.



How Machines Learn?




Futuristic practices and models



What is Machine Learning in the Microsoft Azure cloud?

Azure Machine Learning: Basic workflow


Build models from data and operationalize a machine learning solution



Data Science

The TDSP lifecycle is composed of five major stages that are executed iteratively. These include:

- **1. Business Understanding**
- **2. Data Acquisition and Understanding**
- **3. Modeling**
- **4. Deployment**
- **5. Customer Acceptance**



The 5 questions data science answers

It might surprise you, but there are only five questions that data science answers:

- Is this A or B?
 - Is this weird?
 - How much – or – How many?
 - How is this organized?
 - What should I do next?


Each one of these questions is answered by a separate family of machine learning methods, called algorithms.



Question 1: Is this A or B? uses classification algorithms

Is this A or B?


Classification algorithms



Question 2: Is this weird? uses anomaly detection algorithms

Is this weird?


Anomaly detection algorithms



Question 3: How much? or How many? uses regression algorithms

How much? How many?


Regression algorithms



Question 4: How is this organized? uses clustering algorithms

How is this organized?


Clustering Algorithms



Question 5: What should I do now? uses reinforcement learning algorithms


What should I do now?

Reinforcement Learning Algorithms




Supervised learning

1. Purpose
2. Process: Data set and test set
3. Components
4. Computation and recommendations
5. Verification
6. Selection




Unsupervised learning


1. Purpose
2. Process
3. Components
4. Computation and recommendation



Azure Machine Learning

- Why Azure Machine Learning?
- Getting started
 - Setting up a Microsoft Azure account
 - Setting up a storage account
 - Setting up an Azure Machine Learning workspace
- Advantages





The algorithms clubbed

- Classification: Predict what class case belongs to

Scenarios: Churn analysis, fraud detection, speech recognition

Algorithms: Boosted Decision Tree, Decision Forest, Logistic Regression,...

- Regression: Predict numerical outcomes

Scenarios: Stock prices prediction, sales forecasts, quality control

Algorithms: Bayesian Linear, Linear Regression, Decision Forest

- Clustering: Discover natural groupings of cases

Scenarios: Customer segmentation, pattern recognition

Algorithms: K-means

- Anomaly Detection: Explore unusual patterns

Scenarios: Network intrusion, fraud detection

Algorithms: One class support vector machine ,PCA based anomaly detection

Supervised, Unsupervised and Reinforcement Learning

Supervised Machine Learning

- Data is labelled
- There is an Input variable "X" or set of input variables and an output variable "Y"
$$Y=f(X)$$
- The function is approximated to predict new values of Y given X
- Examples

Regression- Output variable is a real value such as Amount, Height, Weight etc

Classification- Output variable is a category, such as Yes, No, Red, Blue, Yellow etc

Loan_ID	Gender	Married	Dependents	Self_Employed	Income	LoanAmt	Term	CreditHistory	Property_Area	Status
LP001002	Male	No	0	No	\$5,849.00		60	1	Urban	Y
LP001003	Male	Yes	1	No	\$4,583.00	\$128.00	120	1	Rural	N
LP001005	Male	Yes	0	Yes	\$3,000.00	\$66.00	60	1	Urban	Y
LP001006	Male	Yes	2	No	\$2,583.00	\$120.00	60	1	Urban	Y

Unsupervised Machine Learning

- Only X or input variable is known
- The goal for unsupervised learning is to model the underlying structure or distribution in the data in order to learn more about the data.
- There is no correct answers and there is no teacher.
- Algorithms are left on their own to discover and present the interesting structure in the data.
- Examples

Clustering - Customer behaviour grouping
Association - Recommendation model



Customers who viewed this item also viewed these products



Dualit Food XL1500 Processor
\$560

Add to cart



Kenwood kMix Manual Espresso Machine
★★★★★
\$250

Select options



Weber One Touch Gold Premium Charcoal Grill-57cm
\$225

Add to cart




NoMU Salt Pepper and Spice Grinders
\$3

View options

Reinforcement Learning

- Reinforcement learning rewards good behaviour and penalizes bad ones
- The idea is to maximise the gain or reward



Understanding Data, Variables/Features

Understanding The Variables Using a Dataset

Loan_ID	Gender	Married	Dependents	Self_Employed	Income	LoanAmt	Term	CreditHistory	Property_Area	Status
LP001002	Male	No	0	No	\$5,849.00		60	1	Urban	Y
LP001003	Male	Yes	1	No	\$4,583.00	\$128.00	120	1	Rural	N
LP001005	Male	Yes	0	Yes	\$3,000.00	\$66.00	60	1	Urban	Y
LP001006	Male	Yes	2	No	\$2,583.00	\$120.00	60	1	Urban	Y

Types of Variables

- Predictor/ Independent
 - Gender
 - Married
 - Dependents
 - Self_Employed
 - Income
 - Loan Amt
 - Term
 - Credit History
 - Property Area
- Target/Dependent
 - Status

Data Type

- Character/String
 - Gender
 - Married
 - Self-Employed
 - Property Area
 - Status
- Numeric
 - Dependents
 - Income
 - Loan Amt
 - Term
 - Credit History

Category

- Categorical
 - Gender
 - Married
 - Credit History
 - Self-Employed
 - Property Area
 - Status
- Continuous
 - Dependents
 - Income
 - Loan Amt
 - Term

Common Terms in Machine Learning

Mean and Median

- Mean - Average of all the values

$$\begin{aligned}\text{Mean} &= \text{Sum of Salary}/\text{Number of observations} \\ &= 62,800/11 \\ &= \$ 5709.09\end{aligned}$$

- Median - Numerical Middle value of the sorted observations with equal number of observations on both sides,



SALARY
\$ 4,000
\$ 4,400
\$ 5,000
\$ 5,500
\$ 5,700
\$ 5,800
\$ 6,200
\$ 6,400
\$ 6,400
\$ 6,400
\$ 7,000
Sum
\$ 62,800


Mode and Range

- Mode - The value that appears most often in a set of data
6,400
- Range -The difference of highest and lowest values in a sample of observations
 $7000 - 4000 = 3,000$

SALARY
\$ 4,000
\$ 4,400
\$ 5,000
\$ 5,500
\$ 5,700
\$ 5,800
\$ 6,200
\$ 6,400
\$ 6,400
\$ 6,400
\$ 7,000
Sum \$ 62,800

Probability

- Probability is a numerical way of describing how likely something is going to happen.
- Sample Space (S) - Set of possible outcomes that might be observed for an event
 - Dice Sample Space (S) = {1, 2, 3, 4, 5, 6}
- Probability of occurring 3
 - $P(A) = 1/6 = 0.1667$
- Probability of getting an even number from the given sample space
- How many even numbers are there? 2,4,6
- So number even occurrences = 3
- Probability of getting an even number is $P(A) = 3/6 = 0.5$ or 50%



Types of Models

Types of Models


- Classification
- Regression
- Clustering
- Anomaly Detection

Classification

- Classification is identifying to which set of categories a new observation belongs, on the basis of a training set of data containing observations whose category membership is known.
- Binary/ Two-Class Classification — Either/Or, Yes or No type
- Multi-Class Classification — One of the many alternatives

Some examples could be

- Assigning a given email into "spam" or "non-spam" classes Or Primary, Social or Promotional emails
- Will this customer default on loan repayment?
- Will this customer buy my product?




Predicting the value for categorical variable

Regression Analysis

- Regression Analysis is a statistical process for estimating the relationships among variables where the predictor is a continuous variable
- The focus is on the relationship between a dependent variable and one or more independent variables (or 'predictors')
- One of the most common methods used in Machine Learning
- In certain circumstances, it can also be used to infer causal relationships between dependent and independent variables.





Examples

- Predicting the future sale of products
- Computing fair price of the product or service




Casual Relationship?






Clustering or Cluster Analysis

- Clustering is the task of grouping a set of objects in such a way that objects in the same group (called a cluster) are more similar (in some sense or another) to each other than to those in other groups (clusters).
- Unsupervised Learning model
- Customers who make lot of long-distance calls and don't have a job. Who are they?




Cluster Formation



Anomaly Detection

- Anomaly detection (also outlier detection) is the identification of items, events or observations which do not conform to an expected pattern or other items in a dataset.
- Typically the anomalous items will translate to some kind of problem such as
 - Bank fraud
 - Credit Card Fraud
 - Structural defect
 - Medical problems
- Anomalies are also referred to as outliers, novelties, noise, deviations and exceptions.





Getting Started with Azure ML

Walk Through



Previous Section – Basics of Machine Learning


- Why Machine Learning is the Future?
- What is Machine Learning?
- Understanding the data
- Common Machine Learning Terms
- Types of Machine Learning Models

What is Microsoft Azure ML?

Azure Machine Learning

- Cloud based predictive analytics service
- Provides tools to create complete machine learning solution in the cloud
- Quick model creation and deployment using Azure ML Studio
- Allows Models to be deployed as web services
- Provides a large library of Pre-Built Machine learning algorithms and Modules
- Allows for extending your models with custom built R and Python code

Azure ML Studio Overview



Creating Azure ML Account


Demo

Walk Through on Azure ML




Microsoft

Workflow of Azure ML project/experiment



Workflow of Azure ML



Workflow of Azure ML


GET THE DATA

PREPARE THE DATA


Feature Selection

CHOOSE AND APPLY LEARNING ALGORITHMS

TRAIN AND EVALUATE THE MODEL



Azure ML Algorithm CheatSheet



Predicting Categories

Outcome has multiple possibilities? Customer categories etc?

MULTI-CLASS CLASSIFICATION

Fast Training, Linear Model	Multi-Class Logistic Regression
Accuracy, Long Training Times	Multi-Class Neural Network
Accuracy, Fast Training	Multi-Class Decision Forest
Accuracy, Small Memory Footprint	Multi-Class Decision Jungle
Depends on Two-Class	One-V-All Multiclass

Start


TWO-CLASS CLASSIFICATION





Microsoft

Data Processing with Azure ML



Data Input Output to Azure Workspace

Enter Data Manually

Upload a Dataset

Data Format Conversion


Import Data

Unpacked Zipped Dataset

DEMOS

Data Preparation in AML

Sample and Split Data



Dividing your data into multiple subsections of different sizes.

Why to Partition Data

- The input dataset is huge and have too many observations
- Partition the data to create train and test datasets
- Create different bins of data for cross validation of your results
- Create a random sample of observations or a more balanced dataset for a particular column

Demo on Partitioning

Without Stratification

Rating	Number of Observations
Rating 3	20
Rating 5	5
Total	25

Sampling rate 0.4 or 40%



No Stratification

Rating	Number of Observations
Rating 3	10
Rating 5	0
Total	10

Rating	Number of Observations
Rating 3	5
Rating 5	5
Total	10

Rating	Number of Observations
Rating 3	7
Rating 5	3
Total	10

With Stratification

Rating	Number of Observations
Rating 3	20×0.4 
Rating 5	5×0.4
Total	25

Sampling rate 0.4 or 40%



With Stratification

Rating	Number of Observations
Rating 3	8
Rating 5	2
Total	10



Microsoft

Classification in Azure ML

Logistic Regression



What is Logistics Regression?


- Used to predict the probability of an outcome
- Can be binary - Yes/No or Multiple
- Supervised learning method
- Must provide a dataset that already contains the outcomes to train the model.

Understanding the Logistic Regression

$$Y = f(x)$$

$$Y = b_0 + b_1 X$$


$$\text{No of claims} = 18 + b_1(\text{age})$$



Simple Linear Regression

Logistic Regression?

- Outcome is categorical
- What is the probability of this customer buying this product?




Logistic Regression?

Probability needs to satisfy two basic conditions

- Always positive i.e. > 0
- Always less than or equal to 1


$$Y = b_0 + b_1 X \xrightarrow{\text{To make it positive}} \exp(Y) = \exp(b_0 + b_1 X)$$

 To make it less than 1 divide it with a number greater than P

Logistic Regression?

$$\exp(Y) = \exp(b_0 + b_1 X)$$

$$\frac{\exp(Y)}{\exp(Y) + 1} = \frac{\exp(b_0 + b_1 X)}{\exp(b_0 + b_1 X) + 1}$$




$$e^y$$

$$P = \frac{e^y}{e^y + 1} \text{ probability of success}$$

Logistic Regression?

$$P = \frac{e^y}{e^y + 1}$$

probability of success



$$Q = 1 - P = 1 - \frac{e^y}{e^y + 1} = \frac{e^y + 1 - e^y}{e^y + 1} = \frac{1}{e^y + 1}$$

Probability of Failure


Logistic Regression?

$$\frac{P}{1-P} = e^Y$$

$$\text{Log} \frac{P}{1-P} = Y$$

$$\text{Log} \frac{P}{1-P} = b_0 + b_1 X$$

Plotting Logistic Regression?




Problem Statement

- Automate loan eligibility process
- Identify customers whose loan will be approved

Loan_ID	Gender	Married	Dependents	Self_Employed	Income	LoanAmt	Term	CreditHistory	Property_Area	Status
LP001002	Male	No	0	No	\$5,849.00		60	1	Urban	Y
LP001003	Male	Yes	1	No	\$4,583.00	\$128.00	120	1	Rural	N
LP001005	Male	Yes	0	Yes	\$3,000.00	\$66.00	60	1	Urban	Y
LP001006	Male	Yes	2	No	\$2,583.00	\$120.00	60	1	Urban	Y

Demo on Loan Eligibility

Logistic Regression in Azure ML



Properties Project

Two-Class Logistic Regression

Create trainer mode

Single Parameter

Optimization tolerance

1E-07

L1 regularization weight

1

L2 regularization weight

1

Memory size for L-BFGS

20

Random number seed


Allow unknown categorical levels

What are Hyperparameters?




Parameters to Logistic Regression

- Create Trainer Mode
 - Single Parameter - Provide specific set of values
 - Parameter Range - specify multiple values and get the optimum set for given configuration
- Optimization Tolerance - Threshold Value to stop the model iterations on trained dataset
- Memory Size for L-BFGS - Amount of memory to use for next steps and direction
- Random Number Seed - Random integer number that is used for reproducing the same results
- Allow Unknown Categorical Levels - Creates an additional "Unknown" level




Regularization Weight




$$b_0 + b_1 X$$

Under-fit



$$b_0 + b_1 X^2$$

Right



$$b_0 + b_1 X^2 + b_2 X^3$$

Over Fit

What if the effect of such weights is reduced significantly
Or reduced to zero

Regularization Weight


- L2 (Ridge) shrinks all the coefficient by the same proportions but eliminates none
- L1(Lasso) can shrink some coefficients to zero, performing variable selection.
- Both L1 and L2 regularization prevents overfitting by shrinking (imposing a penalty) on the coefficients.
- L2 penalizes one big weight more than many small weights.
- With L2, you tend to end up with many small weights, while with L1, you tend to end up with larger weights, but more zeros.

DEMOS

Prediction Outcome

	Predicted Positives	Predicted Negatives
Actual Positives	True Positives	False Negatives
Actual Negative	False Positives	True Negatives

	Predicted Positives	Predicted Negatives	
Actual Positives	8	2	10
Actual Negative	1	4	5
	9	6	




Accuracy – Proportions of total number of correct results

$$\text{Accuracy} = (8 + 4) / 15 = 0.8 \text{ or } 80\%$$

Prediction Outcome

	Predicted Positives	Predicted Negatives
Actual Positives	True Positives	False Negatives
Actual Negative	False Positives	True Negatives

	Predicted Positives	Predicted Negatives	
Actual Positives	8	2	10
Actual Negative	1	4	5
	9	6	




Precision – Proportion of correct positive results out of all predicted positive results

$$\text{Precision} = 8 / 9 = 0.889 \text{ or } 88.9\%$$

Prediction Outcome

	Predicted Positives	Predicted Negatives
Actual Positives	True Positives	False Negatives
Actual Negative	False Positives	True Negatives

	Predicted Positives	Predicted Negatives	
Actual Positives	8	2	10
Actual Negative	1	4	5
	9	6	




Recall – Proportion of actual positive cases

$$\text{Recall} = 8 / (8 + 2) = 0.8 \text{ or } 80\%$$

Prediction Outcome

	Predicted Positives	Predicted Negatives
Actual Positives	True Positives	False Negatives
Actual Negative	False Positives	True Negatives

	Predicted Positives	Predicted Negatives	
Actual Positives	8	2	10
Actual Negative	1	4	5
	9	6	




F1-Score – Weighted Average (Harmonic Mean) of Precision and Recall

$$\text{F1Score} = 2 * \text{Precision} * \text{Recall} / (\text{Precision} + \text{Recall}) = 0.84$$

Prediction Outcome

	Predicted Positives	Predicted Negatives
Actual Positives	True Positives	False Negatives
Actual Negative	False Positives	True Negatives

	Predicted Positives	Predicted Negatives	
Actual Positives	6	4	10
Actual Negative	0	5	5
	6	6	



In the Previous example

$$\text{Precision} = 6 / 6 = 1 \text{ or } 100\%$$

$$\text{Recall} = 6 / (6 + 4) = 0.6 \text{ or } 60\%$$

$$\text{Average} = 0.8$$

May lead to false interpretation


$$\text{Precision} = 0.889$$
$$\text{Recall} = 0.8$$

$$\text{Average} = 0.84$$

Prediction Outcome

	Predicted Positives	Predicted Negatives
Actual Positives	True Positives	False Negatives
Actual Negative	False Positives	True Negatives

	Predicted Positives	Predicted Negatives	
Actual Positives	6	4	10
Actual Negative	0	5	5
	6	6	



In the first example

$$\text{Precision} = 6 / 6 = 1 \text{ or } 100\%$$

$$\text{Recall} = 6 / (8 + 2) = 0.6 \text{ or } 60\%$$

$$\text{F1Score} = 0.75$$

$$\text{Precision} = 0.889$$

$$\text{Recall} = 0.8$$

$$\text{F1Score} = 0.84$$

AUC ROC



AUC – Area Under the Curve

ROC – Receiver Operating Characteristics


First used during World War II for the analysis of radar signals

Following the attack on Pearl Harbor in 1941, the United States army began new research to increase the prediction of correctly detected Japanese aircraft from their radar signals.

For this purposes they measured the ability of radar receiver operators to make these important distinctions, which was called the Receiver Operating Characteristics

Provides a single number that lets you compare models of different types.

004 - Logistic Regression - Understanding the results - Excel									
File	Home	Insert	Page Layout	Formulas	Data	Review	View	Add-Ins	Team
H17	X	Y	f(x)	0.5					Tell me what you want to do...
3	B	C	D	E	F	G	H	I	J
4	FPR = FP/(FP + TN)								
5	Threshold	0.05				Threshold	0.03		
6	TP	FN				TP	FN		
7	1019	33	TPR	0.96863		1042	10	TPR	0.99049
8	FP	TN	FPR	0.43342		FP	TN	FPR	0.64431
9	3463	4527				5148	2842		
11	Threshold	0.1				Threshold	0.2		
12	TP	FN				TP	FN		
13	913	139	TPR	0.86787		692	360	TPR	0.65779
14	FP	TN	FPR	0.19312		FP	TN	FPR	0.08811
15	1543	6447				704	7286		
17	Threshold	0.4				Threshold	0.5		
18	TP	FN				TP	FN		
19	404	648	TPR	0.38403		317	735	TPR	0.30133
20	FP	TN	FPR	0.03179		FP	TN	FPR	0.0219
21	254	7736				175	7815		
23	Threshold	0.7				Threshold	0.8		
24	TP	FN				TP	FN		
25	179	873	TPR	0.17015		122	930	TPR	0.11597
26	FP	TN	FPR	0.00914		FP	TN	FPR	0.00526
27	73	7917				42	7948		
28	Data	AUC Plot	+			Threshold	0.9		
						TP	FN		
						60	992	TPR	0.05703
						FP	TN	FPR	0.00275
						22	7968		



Impact Analysis

Multiple Model Analysis

- Scenario 1 - Impact of Split Percentage
- Scenario 2 - Impact of Stratification
- Scenario 3 - Low L1 (0.0001) and High L2 (1)
- Scenario 4 - High L1 (1) and Low L2 (0.0001)
- Scenario 5 - High L1 (1) and High L2 (1)
- Scenario 6 - Low L1 (0.0001) and Low L2 (0.0001)

DEMOS

Multiclass Logistic Regression

Wine Quality Prediction



- Fixed and Volatile Acidity
- Citric acid
- Residual sugar
- Chlorides
- Free and Total Sulphur dioxide
- Density
- pH
- Sulphates
- Alcohol Content

P. Cortez, A. Cerdeira, F. Almeida, T. Matos and J. Reis.

Modelling wine preferences by data mining from physicochemical properties. In Decision Support Systems, Elsevier, 47(4):547-553,2009


DEMOS

What is a Decision Tree?

What is Decision Tree?

- Supervised learning method
- Decision support tool that uses a tree-like graph or model of decisions and their possible consequences
- Various variations such as Boosted Decision Tree, Decision Forest, Decision Jungle
- Can be used for categorical as well as continuous variables

Loan ID	Income Level	Credit Score	Employment	Approved?
L1	Medium	Low	Self-Employed	No
L2	High	Low	Self-Employed	Yes
L3	High	High	Salaried	Yes
L4	Medium	Low	Salaried	Yes
L5	Low	High	Salaried	No
L6	Low	Low	Self-Employed	No
L7	High	Low	Salaried	Yes
L8	Medium	Low	Self-Employed	No
L9	High	High	Self-Employed	Yes
L10	Medium	High	Self-Employed	Yes
L11	High	Low	Salaried	Yes
L12	Medium	High	Salaried	Yes
L13	Medium	High	Self-Employed	Yes



LID	IL	CS	ET	Status
L2	High	Low	SE	Yes
L3	High	High	Salaried	Yes
L7	High	Low	Salaried	Yes
L9	High	High	SE	Yes
L11	High	Low	Salaried	Yes


Pure Subset

LID	IL	CS	ET	Status
L5	Low	High	Salaried	No
L6	Low	Low	SE	No
L14	Low	Low	SE	No
L15	Low	High	SE	No

Pure Subset

LID	IL	CS	ET	Status
L1	Medium	Low	SE	No
L4	Medium	Low	Salaried	Yes
L8	Medium	Low	SE	No
L10	Medium	High	SE	Yes
L12	Medium	High	Salaried	Yes
L13	Medium	High	SE	Yes

Split Further




LID	IL	CS	ET	Status
L1	Medium	Low	SE	No
L4	Medium	Low	Salaried	Yes
L8	Medium	Low	SE	No

Split Further

LID	IL	CS	ET	Status
L10	Medium	High	SE	Yes
L12	Medium	High	Salaried	Yes
L13	Medium	High	SE	Yes

Pure Subset




LID	IL	CS	ET	Status
L1	Medium	Low	SE	No
L8	Medium	Low	SE	No

Pure Subset

LID	IL	CS	ET	Status
L4	Medium	Low	Salaried	Yes

Pure Subset

Decision Tree Terms



Ensemble learning

Everyday Ensemble Learning



Decision?

Is this price fair?

Appreciation of price?

Construction Quality?

Location appropriate?

Neighbourhood?



Decision?



Broker or real estate portal to check fair price, price appreciation

Friend or colleague who stays nearby or stayed in the neighbourhood

Inspection by an architect for quality checks and structural defects.

Decision?

Is this price fair?



Appreciation of price?



Construction Quality?



Majority



Weighted Average

Location appropriate?



Neighbourhood?




Ensemble Learning

- All algorithms have errors
- Collective wisdom is higher than the individual intelligence
- Generate a group of base learners and combined result gives higher accuracy
- Different base learners can use different,
 - Parameters
 - Sequence
 - Training sets etc
- Two major Ensemble Learning Methods
 - Bagging
 - Boosting







Bagging


- Various models are built in parallel
- All models vote to give the final prediction



Boosting

- Train the Decision Tree in a sequence
- Learn from the previous tree by focussing on incorrect observations
- Build new model with higher weight for incorrect observations from previous sequence





Two Class Boosted Decision Tree

Bank Telemarketing

- Goal is to predict if the client will subscribe to a product or not

Number of instances — 45, 211

1. Age
2. Job Type
3. Marital Status
4. Education Level
5. Credit Default?
6. Housing Loan?
7. Personal Loan
8. Contacted Type
9. Contacted Month
10. Last Contacted day
11. Contact Duration
12. Campaign Type
13. P-Days
14. Previous
15. P-Outcome
16. Emp-Var-Rate
17. Consumer Price Index
18. Consumer Confidence Index
19. Euribor 3 Month Rate
20. Number of employees
21. Subscribed?


Two Class Boosted Decision Tree?

- Machine learning model based on the boosted decision trees algorithm
- Based on ensemble learning method
- Among the easiest methods to get top performance
- One of the more memory-intensive learners


DEMOS

Bagging

- Various models are built in parallel
- All models vote to give the final prediction




Bagging



Adult Census Data

Problem statement: Predict whether income exceeds \$50K/yr based on census data



1. Age
2. Workclass
3. Fnlwgt
4. Education
5. Education-Num
6. Marital Status
7. Occupation
8. Relationship
9. Race
10. Sex
11. Capital Gains
12. Capital Losses
13. Hours per week
14. Native Country
15. Income

DEMOS

Decision Forest

IRIS Dataset




Ronald Fisher

Dataset Attributes

1. Sepal length in cm
2. Sepal width in cm
3. Petal length in cm
4. Petal width in cm

5. Class:

- Iris Setosa
- Iris Versicolour
- Iris Virginica



Predicted attribute: Class of iris plant


DEMOS

Support Vector Machine


What is SVM?

- Supervised Learning Algorithm
- Can be used for both Regression as well as Classification
- Mostly used for classification
- The observations are separated by a hyperplane in the space

Vectors




Hyperplane




A Hyperplane separates the two classes

Choosing a Hyperplane




Select a hyperplane that separates the two classes

Choosing the right Hyperplane



Margin is maximum distance between the nearest data points and the hyperplane

What comes first?



Identifying the accurate classes comes first before margin calculation.

DEMO

Support Vector Machine

Tune Model Hyperparameters

Models and Parameters

Two-Class Logistic Regression

1

Create trainer mode

Single Parameter

Optimization tolerance

1E-07

L1 regularization weight

1

L2 regularization weight

1

Memory size for L-BFGS

20

Random number seed

Allow unknown categorical levels

Two-Class Boosted Decision...

1

Create trainer mode

Single Parameter

Maximum number of leaves per tree

20

Minimum number of samples per leaf node

10

Learning rate

0.2

Number of trees constructed

100

Random number seed

Allow unknown categorical levels

Two-Class Support Vector M...

1

Create trainer mode

Single Parameter

Number of iterations

1

Lambda

0.001

Normalize features

Project to the unit-sphere

Random number seed

Allow unknown categorical levels

What are Hyperparameters?



Model parameters?

- Decision Trees
 - Maximum number of leaves per tree
 - Minimum number of samples per leaf node
 - Learning rate
 - Number of trees to construct
- Logistic regression
 - Optimization tolerance
 - L1 regularization weight
 - L2 regularization weight
 - Memory size for L-BFGS

Tune Model Hyperparameters?

- Helps in determining the best possible combination of hyperparameters
- Also known as hyperparameter optimization
- Performance metric to measure
 - Accuracy
 - Precision
 - Recall
 - AUC
 - F1Score

Parameter Sweeping Modes

- Random Grid
- Entire Grid
- Random Sweep

What is a Grid?

- Cartesian Product of Parameters
- Parameter 1 → 1, 2, 3
- Parameter 2 → A, B, C, D

		Parameter 1 →		
		1	2	3
← Parameter 2	A	A, 1	A, 2	A, 3
	B	B, 1	B, 2	B, 3
	C	C, 1	C, 2	C, 3
	D	D, 1	D, 2	D, 3

Random Grid

Parameter 1 →

← Parameter 2

	1	2	3
A	A, 1	A, 2	A, 3
B	B, 1	B, 2	B, 3
C	C, 1	C, 2	C, 3
D	D, 1	D, 2	D, 3

Entire Grid

← Parameter 2


Parameter 1 →

	1	2	3
A	A, 1	A, 2	A, 3
B	B, 1	B, 2	B, 3
C	C, 1	C, 2	C, 3
D	D, 1	D, 2	D, 3

Random Sweep

Parameter 1 → Range 1.....4

Parameter 2 → Range A.....D



Iterations

P1, P2
P1, P2
P1, P2
.
. .
.
P1, P2

DEMOS

Tuning Hyper Parameters



Microsoft


Deploying Models



I have a Model. Now What?

- Develop but how to deploy?
- Model language is not supported
- Difficult to deploy in the current architecture
- Tedious Environment set-up
- Many more...

Steps for Azure ML webservices Deployment



DEMOS

Deploying & Creating Webservice




Microsoft

Regression in AML








Regression Analysis

- Statistical process for estimating the relationships among variables
- Relationship between a dependent variable and one or more independent variables (or 'predictors')
- The predictor is a continuous variable
- Can also be used to infer causal relationships between dependent and independent variables.



Casual Relationship?






Linear Regression

Simple Regression

$$Y = B_0 + B_1 X$$


Multivariate linear regression.

$$Y = B_0 + B_1 X_1 + B_2 X_2 + B_3 X_3 + \dots + B_n X_n$$



Simple Linear Regression

Simple Linear Regression



Simple Regression

$$\hat{Y} = b_0 + b_1X$$

Dependent Variable

Independent Variable

Simple Linear Regression

Hrs Studied (X)	Marks (Y)
0	40
2	52
3	53
4	55
4	56
5	72
6	71
6	88
7	56
7	74
8	89
9	67
9	89
5.38	66.31

X - Mean (A)	Y - Mean (B)	A^2	A*B
-5.38	-26.31	28.99	141.66
-3.38	-14.31	11.46	48.43
-2.38	-13.31	5.69	31.73
-1.38	-11.31	1.92	15.66
-1.38	-10.31	1.92	14.27
-0.38	5.69	0.15	-2.19
0.62	4.69	0.38	2.89
0.62	21.69	0.38	13.35
1.62	-10.31	2.61	-16.65
1.62	7.69	2.61	12.43
2.62	22.69	6.84	59.35
3.62	0.69	13.07	2.50
3.62	22.69	13.07	82.04
		89.08	405.46

$$Y = b_0 + b_1 X$$

$$b_1 = \frac{\sum (X - \bar{X})(Y - \bar{Y})}{\sum (X - \bar{X})^2}$$

$$= 405.46 / 89.08$$

$$= 4.55$$


Simple Linear Regression

Hrs Studied (X)	Marks (Y)
0	40
2	52
3	53
4	55
4	56
5	72
6	71
6	88
7	56
7	74
8	89
9	67
9	89
5.38	66.31

$$Y = b_0 + b_1 X$$


$$b_1 = 4.55$$

$$b_0 = 41.8$$



Common Regression Terms


Ordinary Least Square



Minimum

$$\sum_{i=1}^n (y_i - \hat{y}_i)^2$$

Mean Absolute Error



$$\text{MAE} = \frac{1}{n} \sum_{i=1}^n |y_i - \hat{y}_i|$$


Mean absolute error (MAE) is a quantity used to measure how close forecasts or predictions are to the eventual outcomes.

Root Mean Square Error


$$\text{RMSE} = \sqrt{\frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2}$$

- Very commonly used and makes for an excellent general purpose error metric for numerical predictions.
- Compared to the similar Mean Absolute Error, RMSE amplifies and severely punishes large errors.

Relative Absolute Error





$$\sum_{i=1}^n |y_i - \hat{y}_i|$$



$$\sum_{i=1}^n |y_i - \bar{y}_i|$$

Relative Absolute Error



$$\text{RAE} = \frac{\sum_{i=1}^n |y_i - \hat{y}_i|}{\sum_{i=1}^n |y_i - \bar{y}_i|}$$

Business Problem

Build a model to predict the price of the vehicle based on the available historic data


DEMOS

Linear Regression for Automobile Data

Coefficient of Determination
or R Squared

Coefficient of Determination

How much (what %) of variation in Y is described by the variation in X?



R-Square with an Example

Hrs Studied (X)	Marks (Y)
0	40
2	52
3	53
4	55
4	56
5	72
6	71
6	88
7	56
7	74
8	89
9	67
9	89
5.38	66.31
Mean	

$$Y = 41.8 + 4.55X$$

Predicted Marks \hat{Y}
41.80
50.90
55.45
60.00
60.00
64.55
69.10
69.10
73.65
73.65
78.20
82.75
82.75

$(Y - \bar{Y})^2$	$(\hat{Y} - \bar{Y})^2$
692.22	600.74
204.78	237.47
177.16	117.94
127.92	39.82
106.30	39.82
32.38	3.10
22.00	7.78
470.46	7.78
106.30	53.88
59.14	53.88
514.84	141.37
0.48	270.27
514.84	270.27
3028.77	1844.12
SST	SSR


Coefficient of Determination

Sum of Squares Due to Regression

$$SSR = \sum_{i=1}^n (\hat{y}_i - \bar{y}_i)^2$$

Total Sum of Squares

$$SST = \sum_{i=1}^n (y_i - \bar{y}_i)^2$$



Coefficient of Determination

$$\begin{aligned} R^2 &= \text{SSR/SST} = 1844.12/3028.77 \\ &= 0.60886 \end{aligned}$$

Higher the value → Variation in Y is explained by variation in X.

Gradient Descent

Hypothesis


"Proposed explanation made on the basis of limited evidence as a starting point for further investigation"

$$h(x) = b_0 + b_1 x$$

Find out value of b_0 and b_1 such that

$$Y \sim N(h(x))$$


for the given observations



Example of Linear Regression

Hrs Studied	Marks
0	40
2	52
3	53
4	55
4	56
5	72
6	71
6	88
7	56
7	74
8	89
9	67
9	89

Marks



Cost Function

Hypothesis: $h(x) = b_0 + b_1x$

Hrs Studied	Marks	$b_0 = 0; b_1 = 1$ Marks Predicted
0	40	0
2	52	2
3	53	3
4	55	4
4	56	4
5	72	5
6	71	6
6	88	6
7	56	7
7	74	7
8	89	8
9	67	9
9	89	9


$(Y_i - \hat{Y}_i)^2$
1600
2500
2500
2601
2704
4489
4225
6724
2401
4489
6561
3364
6400

$$\frac{1}{2n} \sum_{i=1}^n (y_i - \hat{y}_i)^2$$


b_0	b_1	Cost
0	1	1944.538

Cost Function Plot

b0	b1	cost
0	1	1944.54
0	2	1610.08
0	3	1311.46
0	5	821.77
0	7	475.46
0	8	356.08
0	10	224.85
0	12	237.00
0	14	392.54
0	15	524.08
0	16	691.46
0	17	894.69
0	18	1133.77
0	20	1719.46
0	21	2066.08



Cost Function with b0 and b1



X axis – b_1


Y axis – b_0

$Z = C(b_0, b_1)$






Gradient Descent




Gradient Descent

Cost Function: $C(b_1)$



Gradient Descent



Batch Gradient Descent

$$b_j := b_j - \alpha f(C_i)$$

Does it for
number of examples
number of features
learning rate

Sum of All before taking one step (epoch)
Long time to reach the bottom

Batch Gradient Descent

Batch Vs Stochastic Gradient Descent

X1	X2	...	Xn

$$b_j := b_j - \alpha f(C_i)$$

Does it for
number of examples
number of features
learning rate


Batch Gradient Descent

Randomly shuffle the dataset

Repeat the steps for every example

Modify the coefficient at every step

Stochastic Gradient Descent





DEMO

Online Gradient Descent in Linear Regression

Regression Using Decision Tree?

Decision Tree Terms







LID	IL	CS	ET	Status
L1	Medium	Low	SE	No
L8	Medium	Low	SE	No

Pure Subset


LID	IL	CS	ET	Status
L4	Medium	Low	Salaried	Yes

Pure Subset

Decision Tree Regression





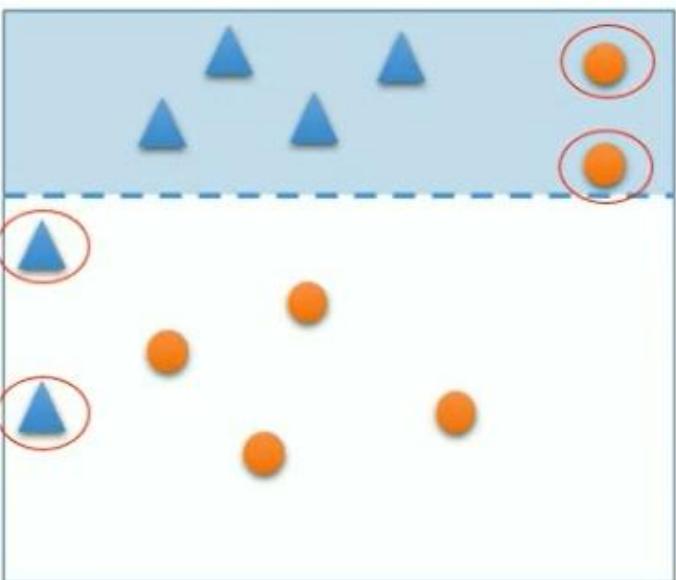



Decision Tree Regression



The output by every region is the mean of that region

Or

Every leaf may have a regression line for the data points within that region



Boosted Decision Tree Regression

- MART gradient boosting algorithm.
- Builds each regression tree in a step-wise fashion
- Predefined loss function to measure the error in each step


DEMO

Boosted Decision Tree Regression

Cluster Analysis

Clustering or Cluster Analysis

- Clustering is the task of grouping a set of objects
- Unsupervised Learning model
- Discovering distinct groups in customer databases
- Used for creating strategies to adopt for certain segments




Examples of Clustering

- Recommendation engines
- Market Segmentation
- Social Network Analysis



Customers who viewed this item also viewed these products




Examples of Clustering


- Medical /Health Science
- Image Segmentation




- Anomaly detection




Euclidean Distance



How Clusters are formed?



https://en.wikipedia.org/wiki/K-means_clustering User: Chire



Good Cluster Analysis

- Observations in the same group share similar characteristics
- Clusters have proportionate number observations

Cluster Initialization

- Random - Random placement of data points into clusters
- First N or Forgy Method - First Data points at Random
- K-Means++ - Default method and an improvement over finding the initial means
- K-Means++ Fast — Optimised for faster clustering
- Evenly
- Use Label Column

DEMO

Clustering using K-Means



Microsoft

Recommendation




AzureML Matchbox Recommender

Recommendation System

- What is a Recommendation System?
- Types of Recommendation Systems
 - Collaborative Filtering
 - Content Based Filtering
- How a recommendation system works?

What is Recommendation System?

"A recommender system or a recommendation system (platform or engine) seeks to predict the "rating" or "preference" that a user would give to an item."




Everyday Recommendations



History of preferences or behaviour



Friends, Family, Colleagues, Professors



Item Features such as Food, taste, product range etc



Recommend




Types of Recommendations Systems

- Collaborative Filtering
- Content Based Filtering
- Hybrid - Combination of the Collaborative and Content Based
- Popularity based - Most bought, most watched, most downloaded, most heard




Collaborative Filtering

- Analyse User behaviour, Activities and preferences
- Recommend based on similarity to other user
- People who agreed in the past will agree in the future, and that they will like similar kinds of items as they liked in the past.




People who buy “x” also buy “y”.

Content based Filtering



Recommend



Business Problem

A travel booking website books thousands of hotel rooms every day.

With such a huge growth, the task at hand is to improve the booking per customer by showing the best hotels as per the taste and preferences of the user.

Sample Set

- Hotel 1 -Good Gym, Small Pool
- Hotel 2 - Amazing Gym, No Pool
- Hotel 3 - Great Gym, Small Pool
- Hotel 4- Basic Gym, Large Pool
- Hotel 5-No Gym, Large and Beautifully designed pool

- John - Prefers Gym over swimming pool
- Kavin - Likes Gym more than pool
- Bill - Needs a pool with basic Gym
- Frans - A pool is a must compared to Gym

Item – Feature Matrix

Hotel	Gym	Pool
Hotel 1	0.8	0.2
Hotel 2	1	0
Hotel 3	0.9	0.1
Hotel 4	0.1	0.9
Hotel 5	0	1

User – Feature Matrix

User	Gym	Pool
John	0.9	0.1
Kavin	0.8	0.2
Bill	0.3	0.7
Frans	0	1


Creating Feature Vector

Hotel	Gym	Pool
Hotel 1	0.8	0.2
Hotel 2	1	0
Hotel 3	0.9	0.1
Hotel 4	0.1	0.9
Hotel 5	0	1



Hotel	Feature Vector
H1	[0.8, 0.2]
H2	[1, 0]
H3	[0.9, 0.1]
H4	[0.1, 0.9]
H5	[0, 1]

User	Gym	Pool
John	0.9	0.1
Kavin	0.8	0.2
Bill	0.3	0.7
Frans	0	1



User	Feature Vector
U1	[0.9, 0.1]
U2	[0.8, 0.2]
U3	[0.3, 0.7]
U4	[0, 1]

Creating Feature Vector

Hotel	Feature Vector	User	Feature Vector
H1	[0.8, 0.2]	U1	[0.9, 0.1]
H2	[1, 0]	U2	[0.8, 0.2]
H3	[0.9, 0.1]	U3	[0.3, 0.7]
H4	[0.1, 0.9]	U4	[0, 1]
H5	[0, 1]		

- Which Hotels should be recommended to John (U1)?
- MAX (Uj * Hi)

MAX (U1*H1, U1*H2, U1*H3, U1*H4, U1*H5)

MAX (0.74, 0.9, 0.82, 0.18, 0.1)

Recommended Hotels in the order of preference

H2, H3, H1

Collaborative Filtering

Workflow of CF, (Wikipedia)

- A user expresses his or her preferences by rating items.
- Ratings can be viewed as an approximate representation of the user's interest in the corresponding domain.
- The system matches this user's ratings against other users' and finds the people with most "similar" tastes.
- With similar users, the system recommends items that the similar users have rated highly but not yet being rated by this user (presumably the absence of rating is often considered as the unfamiliarity of an item)

User	Hotel 1	Hotel 2	Hotel 3	Hotel 4	Hotel 5
John	4	5	?	?	1
Kavin	?	5	4.5	?	1
Bill	2	3	?	5	4
Frans	1	?	?	?	5

Collaborative Filtering

User	Hotel 1	Hotel 2	Hotel 3	Hotel 4	Hotel 5
John	4	5	?	?	1
Kavin	?	5	4.5	?	1
Bill	2	3	?	5	4
Frans	1	?	?	?	5

Matchbox Recommender

Recommendation Data Issues

- Cold Start
- Same Scale Judgement

Cold Start

- Content based filtering constructs user profile before system can recommend
- Collaborative Filtering needs like-minded users for recommendations
- Completely new profile or item
- Can be tackled using Hybrid recommendation

Same Scale judgement

- A user rates all the items on a same scale
- An item is rated on the same scale by all users

Recommender Split

- Fraction of training only users - Fraction of users assigned only to the training dataset. The rows would never be used to test the model.
- Fraction of test user ratings for training - Portion of the user ratings that can be used for training.
- Fraction of cold users - Cold users are users that the system has not previously encountered
- Fraction of cold items - Cold items are items that the system has not previously encountered.
- Fraction of ignored users - Specify the percentage of users that should be ignored.
- Fraction of ignored items - Specify the percentage of items to ignore.
- Remove occasionally produced cold items

DEMOS


Matchbox Recommendation System

What is Matchbox Recommendation System

- Recommendation system developed by Microsoft Research
- Based on a the Hybrid Approach of Content and Collaborative filtering
- Uses Train Matchbox Recommender and Score Matchbox Recommender Modules
- Train Matchbox Recommender reads a dataset of user-item-rating triples and, optionally, some user and item features.

Configuring Train Matchbox Recommender

- User-Item-Rating Data - Prepare the data used for training, containing **user-item-rating** triples
 - First Column – User Identifier
 - Second Column – Item Identifier
 - Third Column – Rating for the user-item pair
- User-Feature – userID and the userfeatures
- Item-Feature – itemID and the Itemfeatures



Parameters to Train Matchbox Recommender

- **Number of Traits**
 - How many traits to learn for each user and item
 - Each feature is associated with a latent “trait” vector
- **Number of recommendation algorithm iterations**
 - how many times the algorithm should process the input data
- **Number of training batches**
 - Number of batches for dividing the data during training

Properties Project

Train Matchbox Recommender

Number of traits

10

Number of recommendation algorithm iterations


5

Number of training batches

4

Score Matchbox Recommender

Input to Score Matchbox Recommender



Types of Prediction

- Predict Ratings
- Item Recommendation
- Related Users
- Related Items

Ratings of Prediction

- Rating predictions by a user
- Input data must contain User and Item
- Does not require any parameter

Item Recommendation

- Users and Items as input
- Uses its knowledge about existing items and users
- Generates a list of items that will appeal to each user

Score Matchbox Recommender

Recommender prediction kind

Item Recommendation

Recommended item selection

From Rated Items (for model evaluation)

Maximum number of items to recommend to a user

5

Minimum size of the recommendation pool for a single user

2

Whether to return the predicted ratings of the items along with the labels

Related Users and Items

- Can be used for "People like you" predictions
- Generate recommendations for users based on items that have already been rated

Understanding the Recommender Result

- For Item Recommendations
- Normalized Discounted Cumulative Gain (NDCG)



Ranking Quality

Search Result

What is machine learning? - Definition from WhatIs.com

[whatis.techtarget.com › Topics › AppDev › Programming](https://whatis.techtarget.com/Topics/AppDev/Programming) ▾

Jun 24, 2017 - Machine learning is the field of study that gives computers the ability to learn without being explicitly programmed.

The 10 Algorithms Machine Learning Engineers Need to Know

www.kdnuggets.com/2016/08/10-algorithms-machine-learning-engineers.html ▾

Aug 8, 2016 - It is no doubt that the sub-field of machine learning / artificial intelligence has increasingly gained more popularity in the past couple of years.

Machine Learning | SAP - SAP.com


<https://www.sap.com/india/trends/machine-learning.html> ▾

Discover how AI, machine learning, and deep learning are powering a new breed of software that uses Big Data to drive radical changes to business.

Machine Learning | edX


<https://www.edx.org/course/machine-learning-columbiasx-csmm-102x-0> ▾

Master the essentials of machine learning and algorithms to help improve learning from data without human intervention.



The highest ranked item should result in highest gain.

Ranking Quality



Discounted Cumulative Gain

Ranked items by the algorithms  I1, I2, I3, I4, I5, I6

$$\begin{aligned} \text{CG} &= 4 + 3 + 4 + 0 + 1 + 2 \\ &= 14 \end{aligned}$$

Gain perceived by the user  4, 3, 4, 0, 1, 2

i	rel_i	$\log_2(i+1)$	$\frac{\text{rel}_i}{\log_2(i+1)}$
1	4	1	4.00
2	3	1.585	1.89
3	4	2	2
4	0	2.322	0
5	1	2.585	0.39
6	2	2.81	0.71
DCG			8.99

Ideal Discounted Cumulative Gain

Ideal Ranking by algorithm



I1, I3, I2, I6, I5, I4

Ideal Gain perceived by the user



4, 4, 3, 2, 1, 0

$$\begin{aligned} \text{CG} &= 4 + 4 + 3 + 2 + 1 + 0 \\ &= 14 \end{aligned}$$

i	rel_i	$\log_2(i+1)$	$\frac{\text{rel}_i}{\log_2(i+1)}$
1	4	1	4.00
2	4	1.585	2.52
3	3	2	1.5
4	2	2.322	0.86
5	1	2.585	0.38
6	0	2.81	0
IDCG			9.27

Normalised DCG

DCG = 8.99

IDCG = 9.27

$$\text{NDCG} = \frac{\text{DCG}}{\text{IDCG}} = \frac{8.99}{9.27} = 0.9697$$

Highly relevant documents are more useful than marginally relevant documents, which are in turn more useful than non-relevant documents

Understanding the Result

- Normalized Discounted Cumulative Gain (NDCG)




Azure AI Stack

Components of Azure Machine Learning

- Azure Machine Learning Workbench
- Azure Machine Learning Experimentation Service
- Azure Machine Learning Model Management Service
- Microsoft Machine Learning Libraries for Apache Spark (MMLSpark Library)
- Visual Studio Code Tools for AI

AZURE MACHINE LEARNING

AZURE MACHINE LEARNING SERVICES



TRAIN & DEPLOY OPTIONS

AZURE



Spark
SQL Server
Virtual machines
GPUs
Container services

ON-PREMISES



SQL Server
Machine Learning Server

EDGE COMPUTING



Azure IoT Edge

Azure Machine Learning Workbench

- Azure Machine Learning Workbench is a desktop application plus command-line tools
- It allows you to manage machine learning solutions through the entire data science life cycle:
 - Data ingestion and preparation
 - Model development and experiment management
 - Model deployment in various target environments

Azure Machine Learning Experimentation Service

- The Experimentation Service handles the execution of machine learning experiments.
- It also supports the Workbench by providing project management, Git integration, access control, roaming, and sharing.
- Through easy configuration, you can execute your experiments across a range of compute environment options:
 - Local native
 - Local Docker container
 - Docker container on a remote VM
 - Scale out Spark cluster in Azure

Azure Machine Learning Model Management Service


- Model Management Service allows data scientists and dev-ops teams to deploy predictive models into a wide variety of environments.
- Model versions and lineage are tracked from training runs to deployments. Models are stored, registered, and managed in the cloud
- Using simple CLI commands, you can containerize your model, scoring scripts and dependencies into Docker images. These images are registered in your own Docker registry hosted in Azure (Azure Container Registry). They can be reliably deployed to the following targets:
 - Local machines
 - On-premises servers
 - The cloud
 - IoT edge devices

Offering from Microsoft on AI

- Besides Azure Machine Learning, there are a wide variety of options in Azure to build, deploy, and manage machine learning models.
 - Microsoft Machine Learning Services in SQL Server
 - Microsoft Machine Learning Server
 - Data Science Virtual Machine
 - Spark MLLib in HDInsight
 - Batch AI Training Service
 - Microsoft Cognitive Toolkit
 - Microsoft Cognitive Services



Cloud AI Stack



Thank you